

Reproducibility Sheet — CS7641 OL Report

Optimization & Uncertainty in Learning, Summer 2026

Timothy Leung

tleung37@gatech.edu gtID 904150400

Georgia Institute of Technology, OMSCS

1. Overleaf project (read-only)

Overleaf project: overleaf.com/project/6a11e4cebb9a5dd9d3751c28

The project is hosted on Georgia Tech Enterprise Overleaf (overleaf.gatech.edu); the read-only link above grants the grader full revision-history access without write permission. Revision timeline is intact from first commit through final submission build.

2. GitHub commit hash (single SHA)

Final commit SHA: bcadfceebf5342f7775caf9aba20f8e86cec3c3e

Repository: github.gatech.edu/tleung37/CS7641_ml_report2 (GT Enterprise, private)

Branch: main

Tag: v1.0-submission

The repository contains every script that produces every figure, every table, and every number in the report. Branch and commit history are intact and available for review.

3. Exact run instructions (standard Linux machine)

3.1 Environment setup

Tested on Ubuntu 22.04 LTS, Python 3.11, R 4.3. Approximate end-to-end runtime: 35–45 minutes on a 4-core CPU laptop (CUDA GPU reduces to ~8 min).

```
# Clone the repository
git clone https://github.gatech.edu/tleung37/CS7641_ml_report2.git
cd CS7641_ml_report2
git checkout bcadfceebf5342f7775caf9aba20f8e86cec3c3e

# Python environment
python3.11 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt

# R environment (for rsanity/ audit; not required for grading main report)
Rscript -e 'install.packages(c("bookdown","yaml","jsonlite","dplyr","ggplot2",
                             "knitr","kableExtra","nnet","caret"),
          repos="https://cloud.r-project.org")'
```

3.2 Commands to reproduce every figure and table

```
# (a) Run the full pipeline end-to-end.
#     Reads configs/config.yaml, writes results/{logs,tables,figures}/
python run_pipeline.py

# (b) Compile the OL report PDF (twice for cross-references).
cd report
pdflatex -interaction=nonstopmode OL_Report.tex
pdflatex -interaction=nonstopmode OL_Report.tex
cd ..

# (c) Rerun the R sanity check (independent audit of Python results).
Rscript -e 'rmarkdown::render("rsanity/sanity_check_OL.Rmd", output_dir="rsanity")'

# (d) Compile this reproducibility sheet.
cd repro
pdflatex -interaction=nonstopmode REPRO_OL.tex
```

3.3 Data paths

`sklearn.datasets.fetch_covtype()` is called by `src/data_loader.py` on first run; the dataset (~75 MB) caches to `.cache/covtype/`. No manual download is required; the script will fetch from UCI if no cache is present.

3.4 Random seeds

Primary seed: `seed = 7641` (course number; same as SL Report), set in `configs/config.yaml` and propagated to `numpy.random.seed`, `random.seed`, `torch.manual_seed`, and CuDNN determinism flags. Multi-seed stability runs (Parts 2 and 3) additionally use seeds `{42, 123, 7}`. Re-running on a clean checkout reproduces every reported number to within floating-point precision on the same Python and `scikit-learn`/`torch` versions.

3.5 Output landing locations

- `results/logs/sl_sgd_baseline.json` — SL PyTorch SGD baseline metrics (the reference point for all OL comparisons).
- `results/logs/{rhc,sa,ga}.json` — Part 1 RO progress and final test metrics per algorithm.
- `results/logs/part2-{optimizer}_seed{s}.json` — Part 2 ablations; one file per optimizer $\times$ seed.
- `results/logs/part3-{regularizer}_seed{s}.json` — Part 3 regularization study; one file per condition $\times$ seed.
- `results/logs/heatmap_alpha_beta1.json` — Adam (α, β_1) sensitivity grid.
- `results/tables/final_metric_table.csv` — the headline leaderboard across all methods.
- `results/figures/*.pdf` — all 10 publication figures.
- `report/OL_Report.pdf` — the compiled main report.
- `rsanity/sanity_check_OL.pdf` — the R audit output.

4. SL Continuity Contract

The OL Report reuses the Covertypes sample, preprocessing pipeline, split, and PyTorch backbone from the SL Report without modification. The table below confirms each component is unchanged.

Table 1: SL $\rightarrow$ OL continuity verification.

Component	SL value	OL value	Changed?
Dataset	Covertypes $n=20k$	Same	No
Seed	7641	7641	No
Split	60/20/20 stratified	Same	No
Preprocessing	StandardScaler (train-only)	Same	No
Architecture	FC(256) $\rightarrow$ FC(128) $\rightarrow$ FC(7)	Same	No
Test set	Stratified 4k rows	Same partition	No

5. Full-data and sampled-data EDA (carried over from SL)

The Covertypes dataset contains 581,012 rows, 54 features (10 continuous + 4 wilderness-area + 40 soil-type binary indicators), and seven cover-type classes.

Table 2: Full Covertypes class distribution. Class-frequency ratio $\approx 103:1$.

Class	Count	Share (%)
1 Spruce/Fir	211,840	36.46
2 Lodgepole Pine	283,301	48.76
3 Ponderosa Pine	35,754	6.15
4 Cottonwood/Willow	2,747	0.47
5 Aspen	9,493	1.63
6 Douglas-fir	17,367	2.99
7 Krummholz	20,510	3.53
Total	581,012	100.00

Table 3: Sampled-data ($n=20,000$) class distribution — same proportions as SL.

Class	Count	Share (%)
1 Spruce/Fir	7,292	36.46
2 Lodgepole Pine	9,751	48.76
3 Ponderosa Pine	1,231	6.16
4 Cottonwood/Willow	95	0.47
5 Aspen	327	1.64
6 Douglas-fir	598	2.99
7 Krummholz	706	3.53
Total	20,000	100.00

A subsequent stratified 60/20/20 split yields a 12,000-row training partition, a 4,000-row validation partition, and a 4,000-row test partition; the test set is the same sealed partition as the SL Report.

6. Exact sampling procedure

Identical to SL Report §6: `StratifiedShuffleSplit(n_splits=1, train_size=20000, random_state=7641)` on the full 581k rows, followed by a second `StratifiedShuffleSplit(test_size=0.20, random_state=7641)` on the 20k sample, then a third `StratifiedShuffleSplit(test_size=0.25, random_state=7641)` on the 16k remainder to produce the internal 12k train / 4k val split.

Standardisation. `StandardScaler` fit on the 12,000-row training partition only; applied to validation and test. The 44 binary indicators pass through unchanged. No test row is seen by any preprocessing step until final evaluation. Implemented in `src/data_loader.py`.

7. Compute accounting

7.1 Gradient-based methods (Parts 2 & 3)

One minibatch backward pass = 1 gradient evaluation (GE) = 1 optimizer update. Batch size 256, $n_{\text{train}}=12,000$: $\lceil 12000/256 \rceil = 47$ GE/epoch. 80 epochs = 3,760 GEs per run (before early stopping). All GE counts logged under `grad_evals` in each JSON log.

7.2 Randomized optimization (Part 1)

One full validation-set evaluation = 1 function evaluation (FE).

Method	Budget	Count rule
RHC	3 restarts $\times$ 400 FE	1 val eval = 1 FE
SA	2,000 FE	1 val eval = 1 FE
GA	50 gen $\times$ pop 20	1 candidate eval = 1 FE

8. RO hyperparameter disclosures

RO trainable layers: final 2 (FC(128) $\rightarrow$ BN $\rightarrow$ ReLU $\rightarrow$ FC(7) + biases). RO trainable parameter count: 1,159 (verified at runtime; $\ll$ 50k cap).

- **RHC**: 3 restarts; perturbation $\mathcal{N}(0, 0.01^2 \mathbf{I})$; restart scale 5σ ; 400 FE/restart.
- **SA**: $T_0=1.0$, geometric cooling $T_t=T_0 \cdot 0.995^t$; Metropolis acceptance; 2,000 FE.
- **GA**: pop 20, tournament ($k=3$), uniform crossover ($p=0.7$), mutation $\mathcal{N}(0, 0.01^2 \mathbf{I})$, top-2 elitism; 50 generations. Progress reported by FE, not generation.

9. Adam-family optimizer disclosures (Part 2)

Optimizer	lr	Other HPs
SGD (no mom.)	0.05	momentum=0
SGD+momentum	0.05	$\mu=0.9$, nesterov=False
Nesterov	0.05	$\mu=0.9$, nesterov=True
Adam	10^{-3}	$\beta_1=0.9$, $\beta_2=0.999$
Adam (no bias-corr.)	10^{-3}	custom; omits $\hat{m}_t$, $\hat{v}_t$
Adam ($\beta_1=0$)	10^{-3}	$\beta_1=0$, $\beta_2=0.999$
AdamW	10^{-3}	$\lambda=10^{-4}$, decoupled WD

Threshold $\ell=0.539$ ($= 0.90 \times$ initial val loss, SGD-no-momentum, seed=42; set from training evidence only, never from test).

Adam without bias correction update rule:

$$\begin{aligned}
 m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t, \\
 v_t &= \beta_2 v_{t-1} + (1 - \beta_2) (g_t)^2, \\
 \theta_t &= \theta_{t-1} - \alpha m_t / (\sqrt{v_t} + \varepsilon).
 \end{aligned}$$

Custom implementation in `src/optimizer.py::AdamNoBiasCorrection`.

10. Regularization disclosures (Part 3)

Standard Adam fixed: $\text{lr}=10^{-3}$, $\beta_1=0.9$, $\beta_2=0.999$, $\varepsilon=10^{-8}$ (no weight decay unless listed).

Condition	Configuration
Adam baseline	No regularization; Adam defaults
L2 10^{-4}	<code>weight_decay=10⁻⁴</code> (coupled L2, not AdamW)
L2 10^{-3}	<code>weight_decay=10⁻³</code>
Early stopping	patience=10 epochs; monitored: val loss; best weights restored
Dropout 0.2	<code>nn.Dropout(0.2)</code> after each hidden layer
Dropout 0.3	<code>nn.Dropout(0.3)</code> after each hidden layer
Label smooth 0.1	<code>CrossEntropyLoss(label_smoothing=0.1)</code>
Input noise 0.05	$x \leftarrow x + \mathcal{N}(0, 0.05^2)$, training only
Best combo	L2(10^{-4}) + early stop(p=10) + label smooth(0.1)

11. R sanity check — method and interpretation

The `rsanity/sanity_check_OL.Rmd` performs two audits:

1. **Baseline agreement.** Trains `nnet` (BFGS, size=64, decay= 10^{-3} , 200 maxit, 8k subsample) on the same seed-7641 stratified split and compares its test macro-F1 to the Python SL SGD baseline loaded from `results/logs/sl_sgd_baseline.json`. $|\Delta_{\text{macro-F1}}| < 0.03$ = confirmed; 0.03–0.10 = expected BFGS-vs-SGD library difference; > 0.10 = investigate. The observed gap is **−0.0371** (R `nnet` = 0.5533 vs. Python SGD = 0.5904), in the expected direction and magnitude, confirming the Python split and preprocessing are faithful. The one-sided negative bias is expected: R uses BFGS on an 8k subsample; Python uses pure SGD on 16k.

2. **Log audit.** Loads every `results/logs/part2_*.json` (21 runs) and `part3_*.json` (27 runs) — 48 files total — and checks: $\text{macro-F1} \in [0, 1]$, GE counts ≥ 0 , no corrupted JSON. All 48 runs pass. Largest seed IQR is 0.035 (stable across seeds). Delta-F1 vs. SL SGD baseline is plotted per regularization condition.

12. Hardware and software versions

- OS: Ubuntu 22.04 LTS (kernel 6.5)
- CPU: 4-core x86_64; no GPU used for submission numbers (CUDA path available but not required)
- Python: 3.11.6
- Key packages: `torch 2.3.1`, `scikit-learn 1.5.0`, `numpy 1.26.4`, `pandas 2.2.2`, `matplotlib 3.9.0`, `seaborn 0.13.2`
- R: 4.3.x; `nnet 7.3`, `caret 6.0`, `bookdown 0.40`, `jsonlite 1.8`, `ggplot2 3.5`
- LaTeX: TeX Live 2025 / TinyTeX with IEEEtran conference class (`IEEEtran.cls v1.8b`)